



Dimitar Zahariev
Technical Trainer

Monitoring Kubernetes Clusters

Партньори:




Следете актуалните обяви за **DevOps**

DEV.BG



Agenda

- Monitoring Kubernetes Clusters
 - Objectives and Challenges
 - Possible Approaches and Solutions
 - Prometheus, Grafana, and Something More
- Demo
- Q&A Session



Monitoring Kubernetes Clusters

Let's explore the possibilities and challenges

Партньори:



Следете актуалните обяви за **DevOps**

DEV.BG



Monitoring vs Observability

- **Monitoring** is the process of collecting, analyzing, and using data to track the performance and health of systems
- Focus Areas
 - **Metrics** - Numerical data representing system performance
 - **Alerts** - Notifications triggered by predefined conditions
 - **Dashboards** - Visual representations of metrics for quick insights
- **Observability** is the ability to understand the internal state of a system by examining its outputs
- Focus Areas
 - **Logs** - Detailed records of events within the system
 - **Metrics** - Quantitative data points for performance analysis
 - **Traces** - Records of the execution path and timings across services
 - **Correlation** - Linking logs, metrics, and traces to form a comprehensive view



Objectives

- **Comprehensive Monitoring**

Achieve end-to-end monitoring of multiple Kubernetes clusters

- **Centralized Metrics**

Aggregate and visualize metrics from all clusters in a single, centralized location

- **Scalability**

Ensure the monitoring solution can scale with the number of clusters and workloads

- **High Availability**

Maintain uptime and reliability of the monitoring system

- **Ease of Use**

Provide intuitive dashboards and alerts for developers and operators



Challenges



- **Scalability Issues**

Single Prometheus instance limitations in a multi-cluster environment

- **Data Storage**

Efficiently storing large volumes of metrics data over time

- **High Availability**

Ensuring the monitoring solution remains available despite failures

- **Multi-tenancy**

Supporting multiple teams or projects with isolated metrics and dashboards

- **Data Aggregation**

Combining metrics from multiple clusters without loss or duplication



Possible Approaches and Solutions

- **Single Prometheus per Cluster**

- Description - Deploy a Prometheus instance in each Kubernetes cluster
- Pros - Simplicity, isolated failure domains
- Cons - Difficult to centralize data, scalability issues, higher maintenance overhead

- **Prometheus + [Thanos | Cortex]**

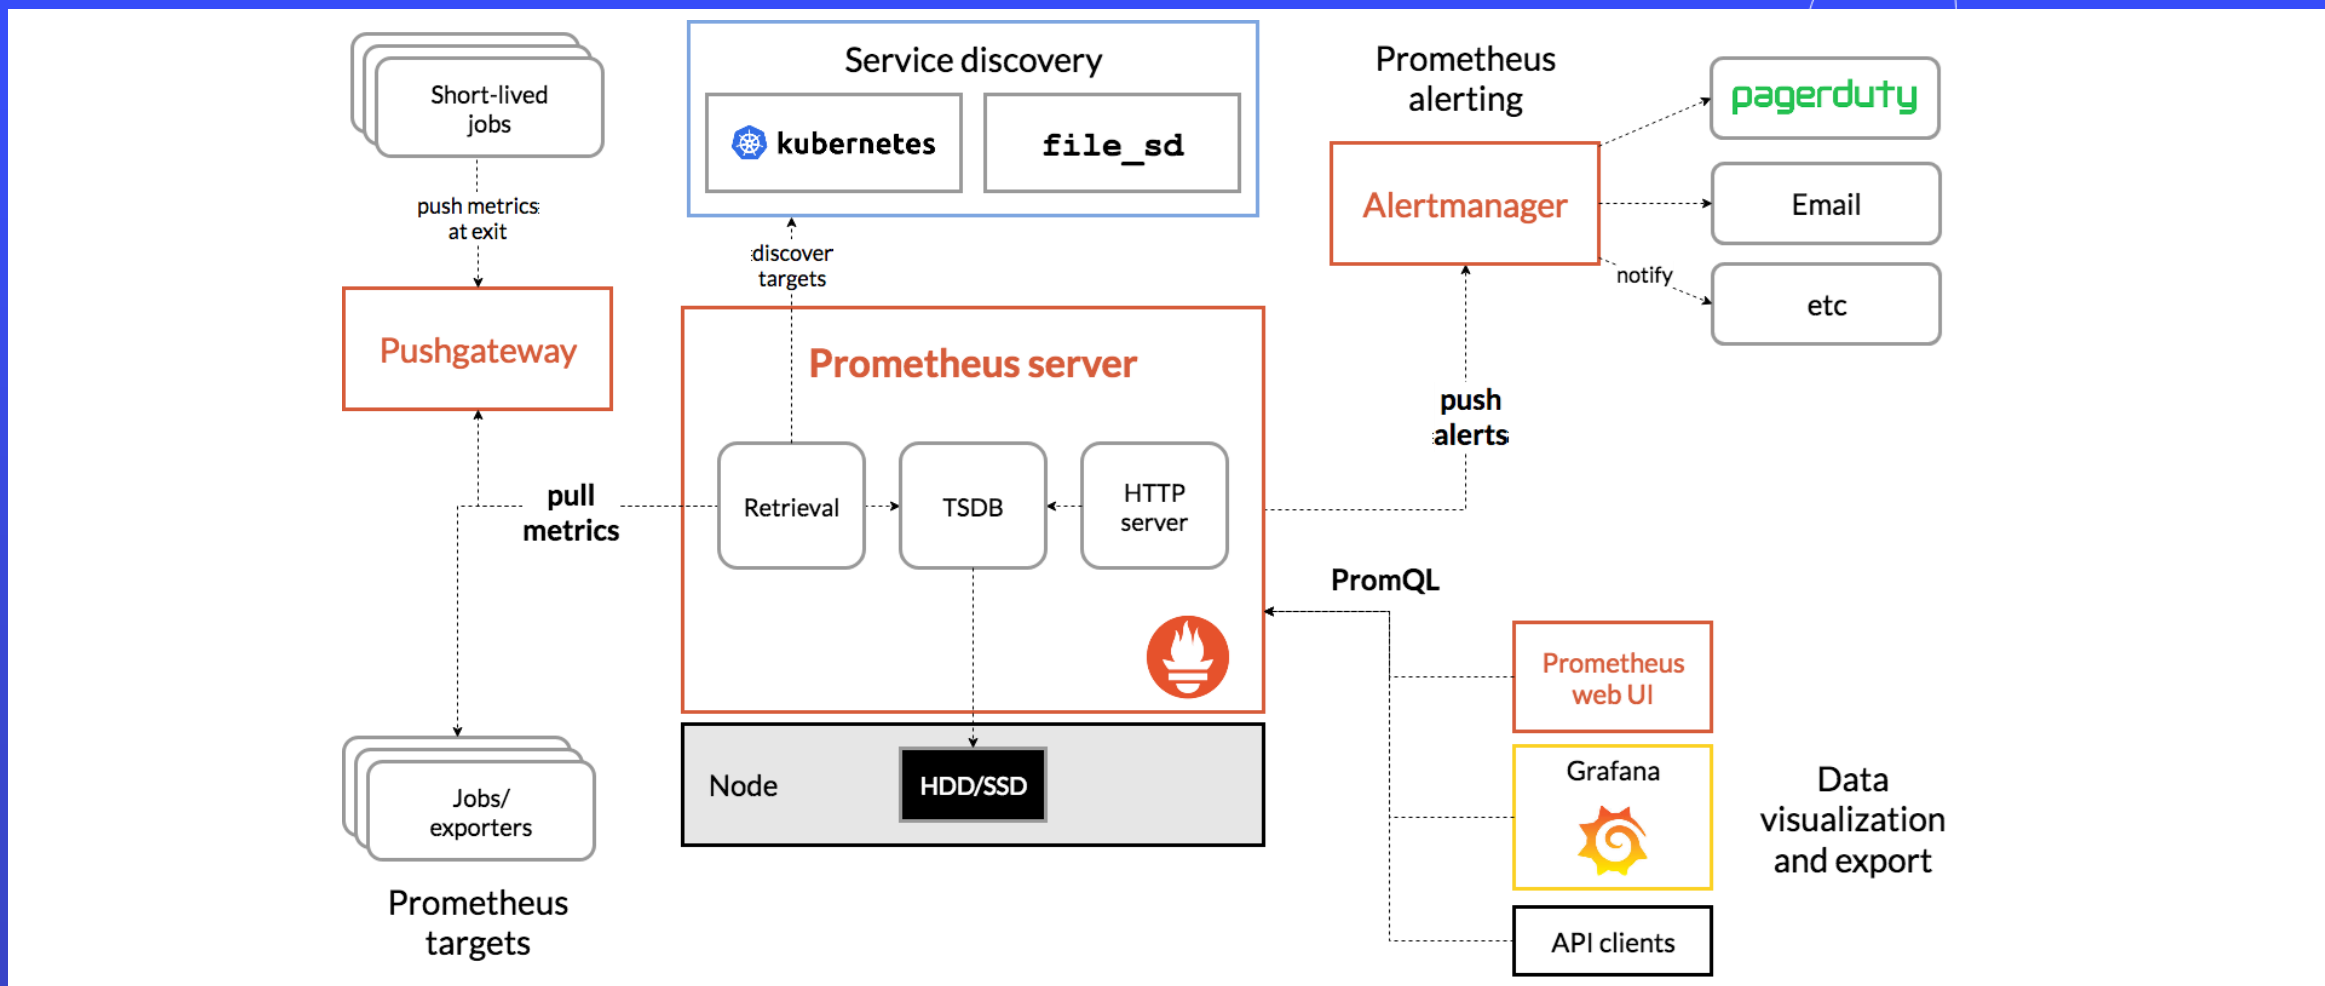
- Description - Use Thanos or Cortex to aggregate/centralize metrics from multiple Prometheus instances
- Pros - Long-term storage, global querying, horizontal scalability, high availability, and multi-tenancy
- Cons - Increased complexity, dependency on external components, requires careful planning and configuration



Prometheus

- An open-source systems monitoring and alerting toolkit
- Main Components
 - **Prometheus Server** - scrapes and stores metrics
 - **Alertmanager** - handles alerts
 - **Exporters** - collect metrics from various sources
- Key Features
 - **Scraping** - collects metrics from endpoints at specified intervals
 - **PromQL** - powerful query language for metric analysis
- Limitations
 - Limited scalability in a multi-cluster environment
 - Challenges in long-term storage and high availability

Prometheus Architecture





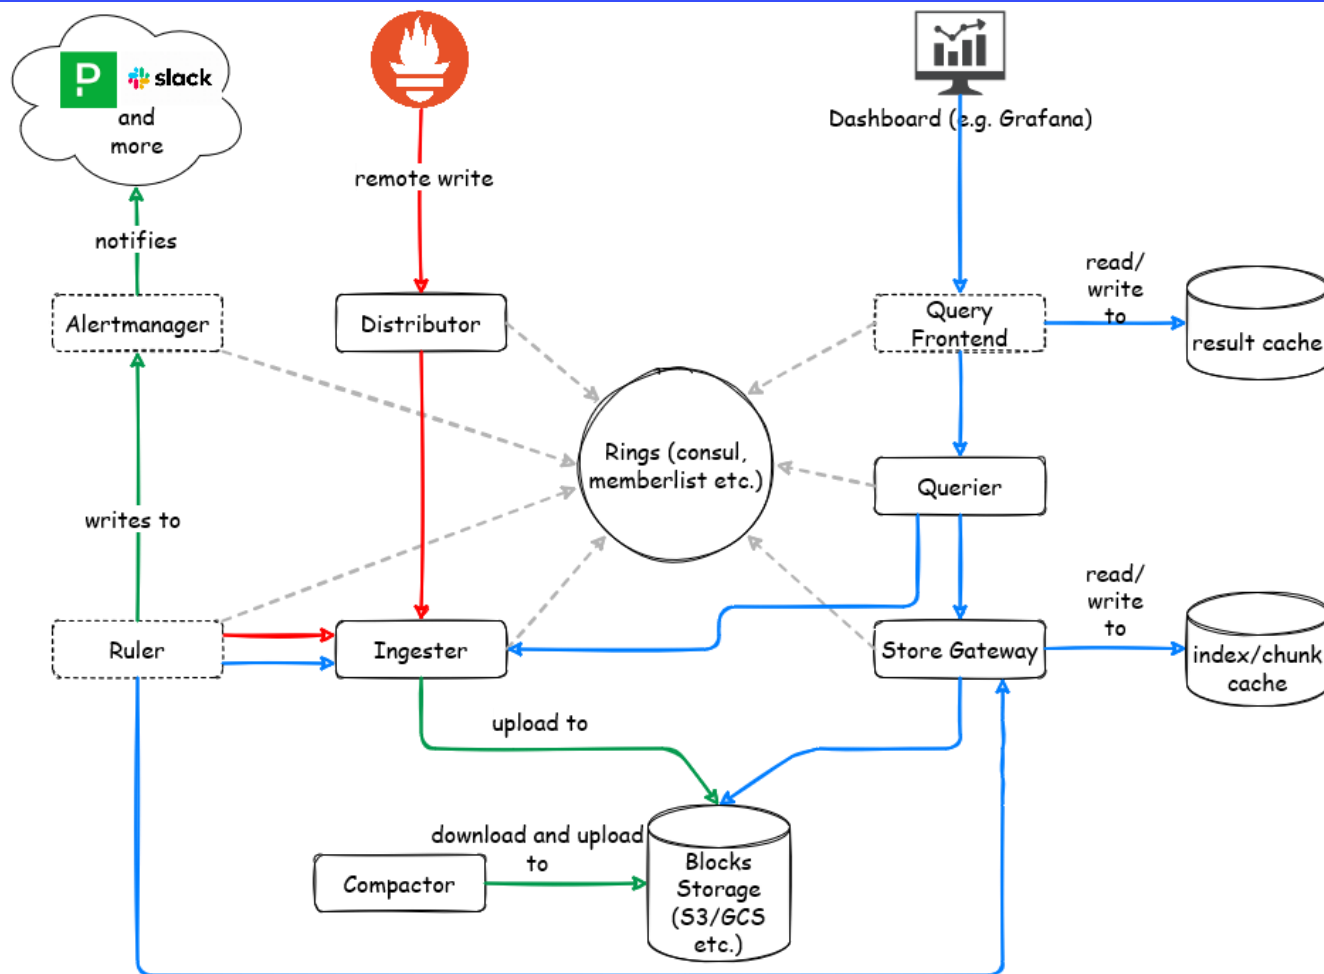
Grafana


- An open-source platform for visualization, monitoring, and observability
- Main Components
 - Dashboards - create and share visual representations of metrics
 - Data Sources - integrate with Prometheus, Cortex, and other data sources
 - Alerting - manage and visualize alerts
- Key Features
 - Flexible and customizable dashboards
 - Wide range of plugins and integrations
- Limitations
 - Dependent on the underlying data source for performance and availability



- Horizontally scalable, highly available, multi-tenant, long term storage for Prometheus
- Main Components
 - **Distributor** - Receive, validate, and route samples to ingesters
 - **Ingestor** - Temporarily store incoming series before writing them to a long-term backend storage
 - **Querier** - Fetch series samples both from the ingesters and long-term storage
 - **Store Gateway** - Responsible for querying series from blocks stored in the long-term storage
- Key Features
 - Horizontal Scalability - Easily scale out by adding more nodes
 - Multi-tenancy - Isolate data for different teams or projects
 - High Availability - Redundant components ensure reliability
 - Long-term Storage - Supports S3, GCS, Swift, and Microsoft Azure for storing metric data

Cortex Architecture





Demo

Let's see it in action

Партньори:

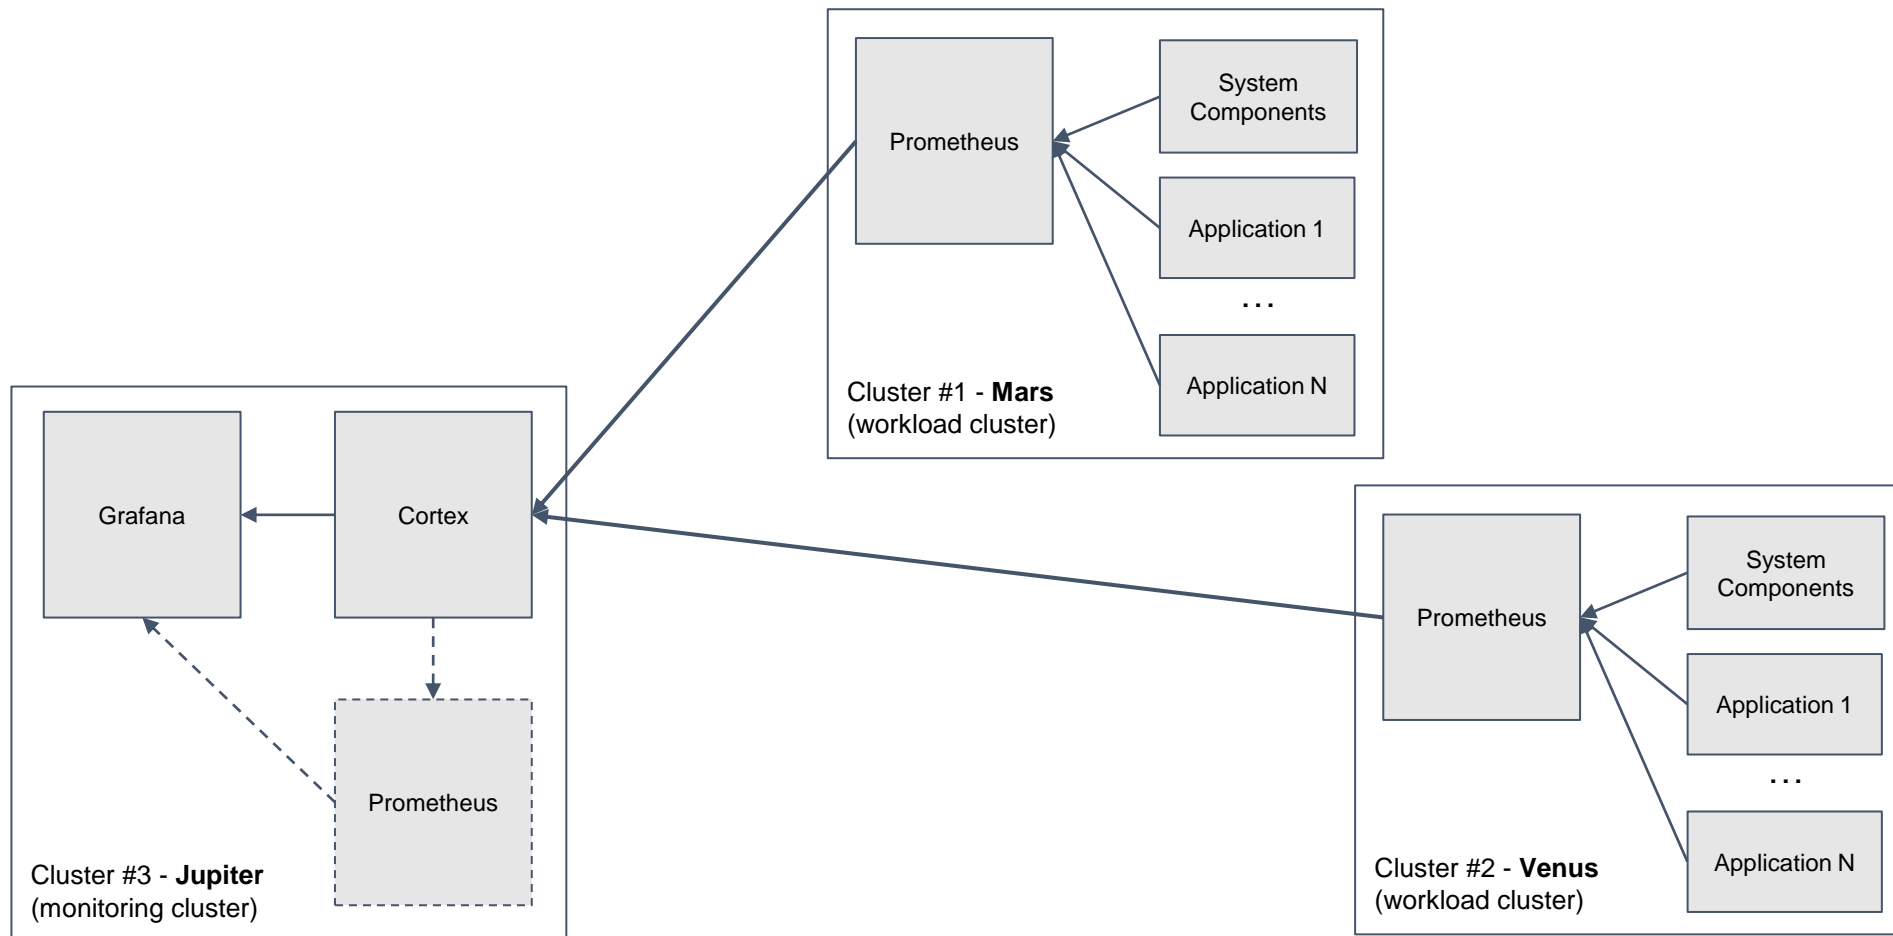



Следете актуалните обяви за **DevOps**

DEV.BG



Our Scenario





Q&A Session

You ask, I answer (if I can)

Партньори:



Следете актуалните обяви за **DevOps**

DEV.BG

Thank you!

Contacts:

 <https://www.linkedin.com/in/dzahariev/>

 <https://github.com/shekeriev>

СЛЕДВАЩО СЪБИТИЕ



Check regularly on **dev.bg** or subscribe to the monthly bulletin

Партньори:



Следете актуалните обяви за **DevOps**

DEV.BG